



# The Generative Intelligence Lab @ FAU

## Comparative Evaluation of Prompt Engineering Techniques for Requirement Generation

---

Exploring the systematic combination of prompt engineering techniques with convergence validation in self-reflective prompting.

- **Authors:** Chris Simmons
  - **Academic Supervisor:** [Dr. Fernando Koch](#)
  - **GitHub Repository:** [github.com/csimmon4-fau/prompt-eng](https://github.com/csimmon4-fau/prompt-eng)
- 

### Research Question

How can we systematically evaluate prompt engineering techniques while leveraging convergence testing in self-reflective prompting to improve requirements quality?

---

### Arguments

#### What is already known about this topic

- **Zero-shot prompting** provides basic requirements elicitation but lacks precision.
- **Few-shot prompting** improves accuracy through examples but needs careful curation.
- **Self-reflective prompting** enables iterative improvement but requires validation.
- **Requirements quality assessment** needs objective measures.
- **Convergence testing** can help validate the stability of iterative processes.

#### What this research is exploring

##### 1. Zero-Shot Baseline

- Basic requirements elicitation capabilities.
- Understanding limitations without context.
- Establishing baseline metrics.

##### 2. Few-Shot Enhancement

- Adding domain-specific examples.
- Improving requirement specificity.
- Measuring contextual impact.

3. Self-Reflective with Convergence

- Implementing iterative refinement.
- Using convergence testing to validate stability.
- Measuring when requirements stabilize.
- Applying similarity thresholds.

Implications for practice

- Provides a systematic approach to requirements elicitation.
- Enables objective measurement of requirement quality.
- Validates requirement stability through convergence testing in self-reflection.
- Establishes a repeatable methodology for requirement refinement.

---

Use Case & Justification

This study focuses on automating requirement analysis for a local, privacy-preserving LLM-based redaction tool. The tool is designed to accurately and efficiently redact names, emails, and other sensitive information from meeting transcripts while ensuring compliance with privacy and security standards.

Why This Problem?

- **Privacy regulations require accurate redaction of sensitive data** in meeting transcripts and other sensitive documents.
- **Manual redaction is error-prone and inefficient**, making automation critical.
- **LLMs offer high accuracy but must be privacy-preserving**, ensuring no sensitive data leaks.
- **Self-reflective prompting can iteratively refine requirement quality** to ensure security, scalability, and efficiency constraints are met.

---

Research Method

Our methodology consists of three phases:

1. **Baseline Evaluation:** Establishing zero-shot and few-shot performance.
2. **Self-Reflective Experimentation:** Iterative refinement using self-reflective prompting.
3. **Performance Measurement:** Evaluating clarity, specificity, and stability over iterations.

---

Metric Definitions

Clarity

Clarity is measured using the **Flesch Reading Ease score**, which evaluates sentence complexity and word difficulty. A higher score indicates the text is easier to read, while a lower score suggests greater complexity.

Specificity

Specificity is evaluated by counting the occurrences of strong requirement-defining words, such as "must," "shall," "exactly," "minimum," and "threshold." Higher specificity scores indicate greater precision and explicit

constraints in the generated requirements.

Effectiveness

Effectiveness is assessed by detecting action-oriented and goal-driven words, such as "ensure," "optimize," "enhance," "reduce," and "automate." This metric captures how well the requirement describes an outcome or functional objective rather than being vague or generic.

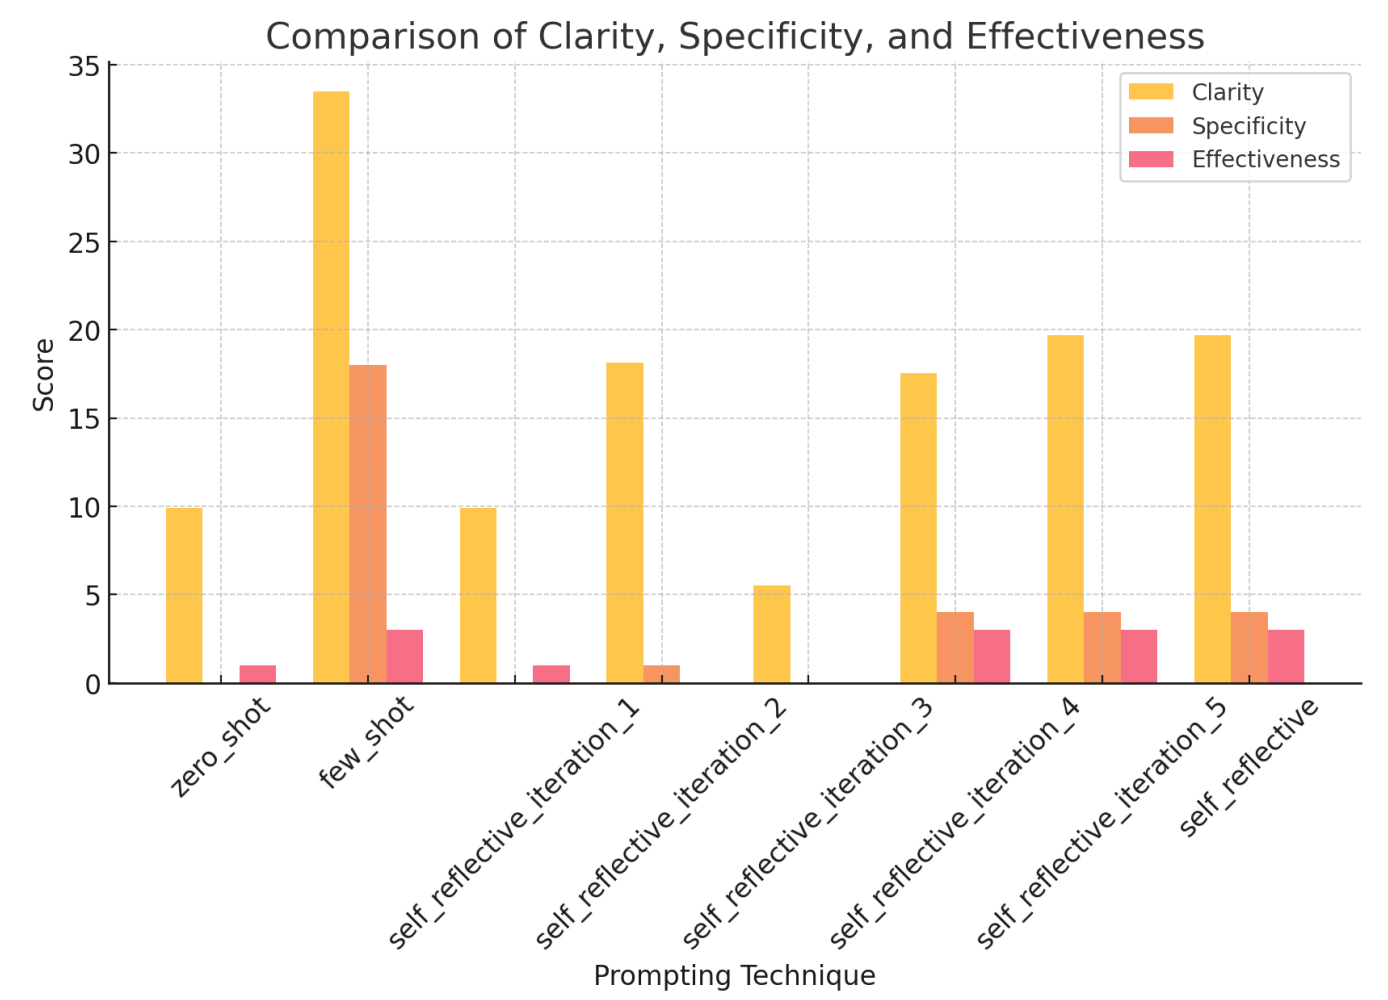
Results & Findings

Comparison of Prompting Techniques

The following table presents the average clarity, specificity, effectiveness scores and run time for each prompting technique.

Prompting Technique	Clarity	Specificity	Effectiveness	Run Time (seconds)
zero_shot	9.89	0.00	1.00	33.902
few_shot	33.51	18.00	3.00	48.103
self_reflective_iteration_1	9.89	0.00	1.00	23.774
self_reflective_iteration_2	18.15	1.00	0.00	23.676
self_reflective_iteration_3	5.53	0.00	0.00	31.329
self_reflective_iteration_4	17.54	4.00	3.00	31.622
self_reflective_iteration_5	19.67	4.00	3.00	36.153
self_reflective - Best Prompt	19.67	4.00	3.00	119.158

Visualization of Metric Trends



**Figure 1: Comparative Analysis of Prompt Engineering Techniques.** This visualization presents how the three prompting methods performed across our key quality metrics.

**Analysis and Findings:**


The data reveals several important patterns:

- Few-shot prompting demonstrates superior performance** in clarity (33.51) and specificity (18), while matching self-reflective prompting in effectiveness (3). This suggests that providing examples is the most effective approach for generating high-quality requirements.
- Self-reflective prompting shows moderate improvement** over zero-shot, with final clarity (19.67) and specificity (4) scores that are better than zero-shot but don't match few-shot performance. The effectiveness score (3) equals few-shot, indicating that iteration can achieve good action-orientation.
- Zero-shot prompting consistently underperforms** with the lowest scores in clarity (9.89), specificity (0), and effectiveness (1), confirming that providing no context or examples produces vague, general requirements.
- Self-reflective iteration is non-linear**, with fluctuating clarity (ranging from 5.53 to 19.67) and inconsistent specificity/effectiveness scores across iterations. This suggests that self-reflection doesn't guarantee steady improvement and may require careful guidance.
- Effectiveness appears to have a ceiling effect** at score 3, achieved by both few-shot and self-reflective methods, indicating that basic action-orientation can be achieved through different

approaches.

These findings suggest that for optimal requirement generation, few-shot prompting provides the most reliable results, while self-reflective prompting may be beneficial as a secondary refinement technique rather than a primary approach.

Visualizations of Metric Trends



**Figure 1: Comparative Analysis of Prompt Engineering Techniques.** This visualization presents a side-by-side comparison of how the three prompting methods (zero-shot, few-shot, and self-reflective) performed across our three key quality metrics. The x-axis groups the methods, while the y-axis shows the metric values for clarity (measured by Flesch Reading Ease), specificity (count of requirement-defining keywords), and effectiveness (frequency of action-oriented language).

Analysis and Findings:

The visualization reveals several important patterns:

- 1. **Few-shot prompting demonstrates clear superiority** across all metrics, with particularly dramatic differences in specificity (18) compared to self-reflective (4) and zero-shot (0) approaches. This indicates that providing examples has the strongest positive impact on requirement quality.
- 2. **Clarity scores follow a consistent pattern** with few-shot (33.51) leading, followed by self-reflective (19.67), and zero-shot (9.89) showing the lowest readability. This progressive improvement suggests that both examples and iteration contribute to producing more readable requirements.
- 3. **Effectiveness shows a threshold effect** rather than a gradient, with both few-shot and self-reflective methods achieving identical scores (3) compared to zero-shot (1). This suggests that once a certain level of guidance is provided (either through examples or reflection), action-oriented language reaches a similar level.

4. **The performance gap between techniques is non-uniform** across metrics, with specificity showing the most dramatic differences. This indicates that the choice of prompting technique may matter more for certain quality aspects than others.

These visual findings reinforce our quantitative analysis and highlight the importance of selecting appropriate prompting techniques based on which requirement quality attributes are most critical for a particular application.

Sample Requirements Generated by Different Techniques

Zero-Shot Example

"The system must implement data redaction capabilities."

Few-Shot Example

"The system must implement data redaction with 99.9% accuracy for personally identifiable information (PII) including but not limited to: names, addresses, phone numbers, and government identification numbers."

Self-Reflective (Final Iteration) Example

"The system shall provide comprehensive data redaction capabilities with a minimum accuracy threshold of 95% for all PII categories, with detailed logging of redaction decisions."

Cost-Benefit Considerations

Technique	Relative Token Usage	Implementation Complexity	Quality Improvement
Zero-Shot	Low (1×)	Minimal	Baseline
Few-Shot	Medium (2-3×)	Low	Substantial (+239% clarity)
Self-Reflective	High (5×)	Medium	Moderate (+99% clarity)

Organizations should consider these tradeoffs when selecting a prompting strategy based on their specific requirements and constraints.

Future Research

1 Hybrid Prompting Strategies for Balancing Specificity & Effectiveness

- The study found that few-shot prompting performs best in clarity, specificity, and effectiveness.
- Self-reflective iterations improve stability but weaken specificity and effectiveness.
- Future research should explore how structured examples can be combined with self-reflective techniques to maintain both consistency and precision in requirements generation.

2 Machine Learning-Based Metrics for Specificity & Effectiveness

- Current evaluation relies on keyword-based scoring.

- Future research could implement ML classifiers or LLM-based assessment models to evaluate requirement quality beyond keyword occurrences.
- This would allow for context-aware scoring that adapts to different domains and requirement structures.

### 3 User-Centric Validation of Prompt Engineering Results

- The current study relies on algorithmic evaluation metrics.
- Future research could introduce human expert assessments to validate whether refined requirements align with stakeholder expectations.
- Crowdsourced evaluations or comparative studies could assess how end-users perceive clarity, specificity, and effectiveness in generated requirements.

---

#### Limitations of the Current Study

1. **Single Model Testing:** Results are specific to the llama3.2 model and may vary with other LLMs
2. **Metric Subjectivity:** While we attempted to quantify quality through objective metrics, some aspects of requirement quality remain inherently subjective
3. **Domain Specificity:** Findings may be particularly applicable to security requirements but could differ for other requirement domains

---

## Implementing Effective Prompt Engineering for Security Requirements

Based on our findings, we recommend the following practical approach:

1. **Start with Few-Shot Examples:** Begin with 3-5 high-quality examples that represent your desired output format and specificity level
2. **Select Domain-Relevant Examples:** Use examples that match your specific security domain (e.g., authentication, data privacy)
3. **Implement Iterative Refinement:** For critical requirements, apply 2-3 rounds of self-reflection to enhance initial outputs
4. **Verify Against Standards:** Cross-check generated requirements against relevant security standards (NIST, ISO 27001, etc.)

#### Connections to Other Disciplines

Our findings on prompt engineering effectiveness have implications for:

1. **Requirements Engineering:** Traditional requirement elicitation might be enhanced by LLM-assisted processes
2. **Security Compliance:** Automated generation of compliance-ready security requirements could accelerate certification processes
3. **Software Testing:** Similar prompting techniques might be applicable for generating comprehensive test cases for security features

#### Acknowledgments

This research was conducted with assistance from generative AI tools, specifically Claude Sonnet 3.7, which was used for editing, formatting, and refining the manuscript. All experimental design, data collection, analysis, and conclusions were human-directed. The use of AI tools aligns with the research focus on prompt engineering techniques and provides a practical demonstration of the capabilities discussed in this paper.